

SSH Attack and Defense

Custom Project Experiment

Peyton Robinson

10/11/2025

Table of contents

Table of contents	2
Summary.....	3
Overview	4
Attack and Defense Timeline.....	5
Timeline Overview	5
Phase 1: Initial Attack	6
Step 1: Deploy Kali Linux virtual attack machine	6
Step 2: Check port 22 availability with Nmap	6
Step 3: Brute force port 22 with Metasploit	8
Phase 2: Defense	10
Step 1: setup Iptable rules	10
Step 2: Add authentication logs into Splunk	11
Step 3: Create Splunk alerts for SSH login attempts	13
Phase 3: Test Attack	13
Step 1: Brute force port 22	13
Step 2: View port 22 with Nmap after attack	13
Step 2: Monitor SSH login alerts in Splunk	14
Outcome	15
Lessons Learned.....	15
Recommendations.....	15
References	17

Summary

Given that SSH remains a critical and popular attack vector, this custom project was executed in a home lab environment across three phases: Initial Attack, Defense Implementation, and Test Attack.

The project used the following tools and technologies:

1. Linux OS target and attack machine
2. Splunk Enterprise
3. Iptables
4. Nmap and Metasploit

The timeline includes supporting screenshots, walking through the initiation of attacks and the subsequent construction of the defenses.

Through this experiment, I gained valuable practical experience on how to leverage Splunk and Iptables to effectively mitigate SSH brute force attacks. I concluded that for future hardening, disabling root access and implementing Multi-Factor Authentication (MFA) would further strengthen the defense posture. For future projects, I plan to incorporate more AI tools and prioritize video tutorials over traditional articles for research.

Overview

The SSH protocol is a popular attack vector, and misconfigurations can lead to malicious connections. Once a connection is established for the attacker, a plethora of other attacks could take place: privilege escalation, lateral movement, and data exfiltration to name a few.

I created this custom attack and defense project to learn about Splunk and firewall defenses against brute force attacks on the SSH protocol.

The relevant components are the following:

1. Linux Host device
2. Kali Linux attack virtual machine
3. Splunk Enterprise
4. Iptables
5. Nmap
6. Metasploit

The attack is only to attempt to establish a SSH connection into the host computer on port 22. The attack does not go any further, whether getting into the host computer or not.

Defense is only to rate-limit SSH logins from the attacker IP and set up alerts in Splunk.

Attack and Defense Timeline

Timeline Overview

The timeline below is a summary of the project's steps.

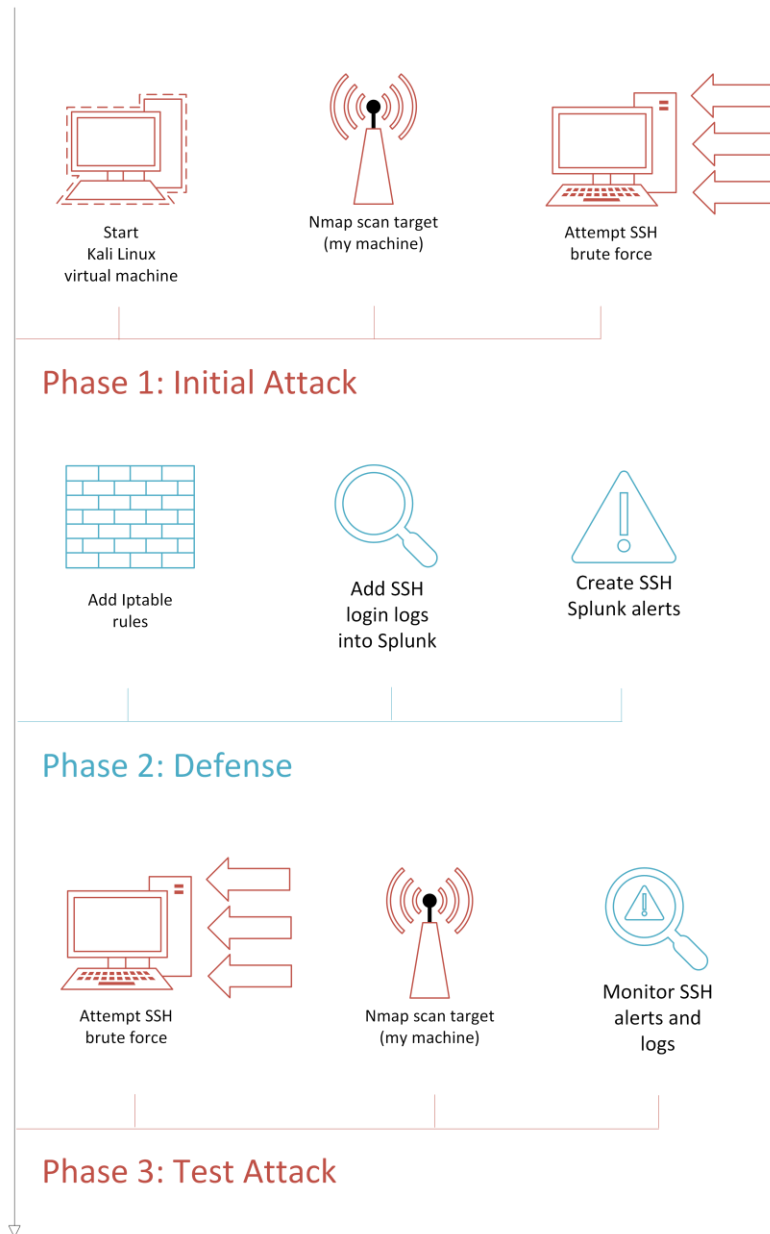


Figure 1 - Microsoft Visio timeline overview of events and phases of attack and defense

Phase 1: Initial Attack

Step 1: Deploy Kali Linux virtual attack machine

To initiate the attack, I started my Kali Linux virtual machine on VirtualBox. This machine was already configured and updated from previous home lab experiments.

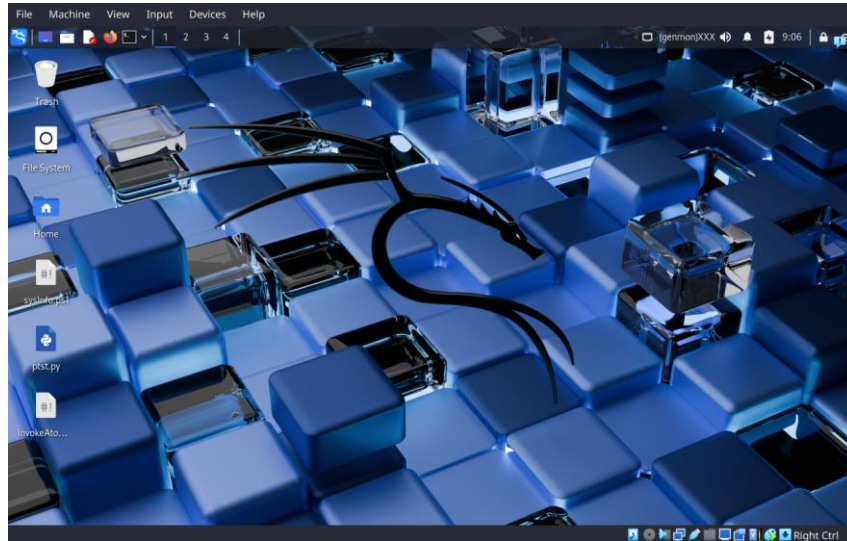


Figure 2 – My Kali Linux home screen.

Step 2: Check port 22 availability with Nmap

Before using Nmap, I installed and enabled SSH on the target machine, which is my personal computer. I used Nmap in the terminal to gather information about port 22 on the target. I also requested extensive machine details by passing the -A flag, and Nmap identified the host as a virtual machine. This misidentification occurred because I bridged the network adapter to the virtual machine, causing Nmap to effectively scan the host's interface.

```
File Machine View Input Devices Help
mb@kali: ~
Session Actions Edit View Help
Nmap done: 1 IP address (1 host up) scanned in 17.56 seconds

(mb@kali)~$ nmap -T2 -A -p 22
Starting Nmap 7.95 ( https://nmap.org ) at 2025-10-09 09:08 EDT
Nmap scan report for [redacted]
Host is up (0.00048s latency).

PORT      STATE SERVICE VERSION
22/tcp    open  ssh      OpenSSH 9.6p1 Ubuntu 3ubuntu13.14 (Ubuntu Linux; protocol 2.0)
|_ ssh-hostkey:
|_ 256 [redacted] (ECDSA)
|_ 256 [redacted] (ED25519)
Warning: OSScan results may be unreliable because we could not find at least 1 open and 1 closed port
Device type: VoIP adapter|bridge|general purpose|printer
Running (JUST GUESSING): AT&T embedded (98%), Oracle Virtualbox (94%), Slirp (94%), QEMU (94%), Konica Minolta embedded (89%), GNU Hurd (88%)
OS CPE: cpe:/a:oracle:vm_virtualbox cpe:/a:danny_gasparovski:slirp cpe:/a:qemu:qemu cpe:/h:konicaminolta:7035 cpe:/o:gnu:hurd
Aggressive OS guesses: AT&T BGW210 voice gateway (98%), Oracle Virtualbox Slirp NAT bridge (94%), QEMU user mode network gateway (94%), Konica Minolta 7035 printer (89%), GNU Hurd 0.3 (88%)
No exact OS matches for host (test conditions non-ideal).
Network Distance: 2 hops
Service Info: OS: Linux; CPE: cpe:/o:linux:linux_kernel

TRACEROUTE (using port 22/tcp)
HOP RTT ADDRESS
1 0.26 ms [redacted]
2 0.51 ms [redacted]

OS and Service detection performed. Please report any incorrect results at https://nmap.org/submit/ .
Nmap done: 1 IP address (1 host up) scanned in 28.63 seconds

(mb@kali)~$
```

Figure 3 – Nmap scan result of the target machine

Metasploit is a common penetration testing tool used to automate attacks on devices. I utilized its capability to brute-force SSH connections

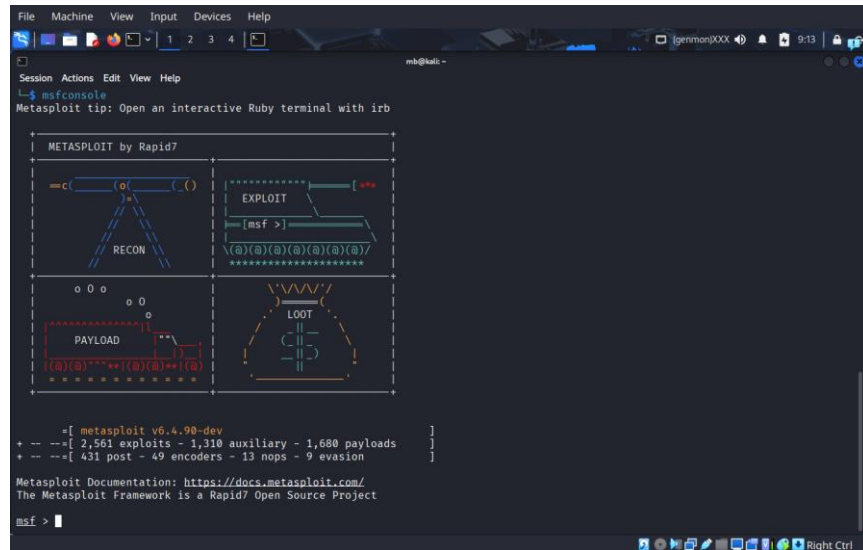
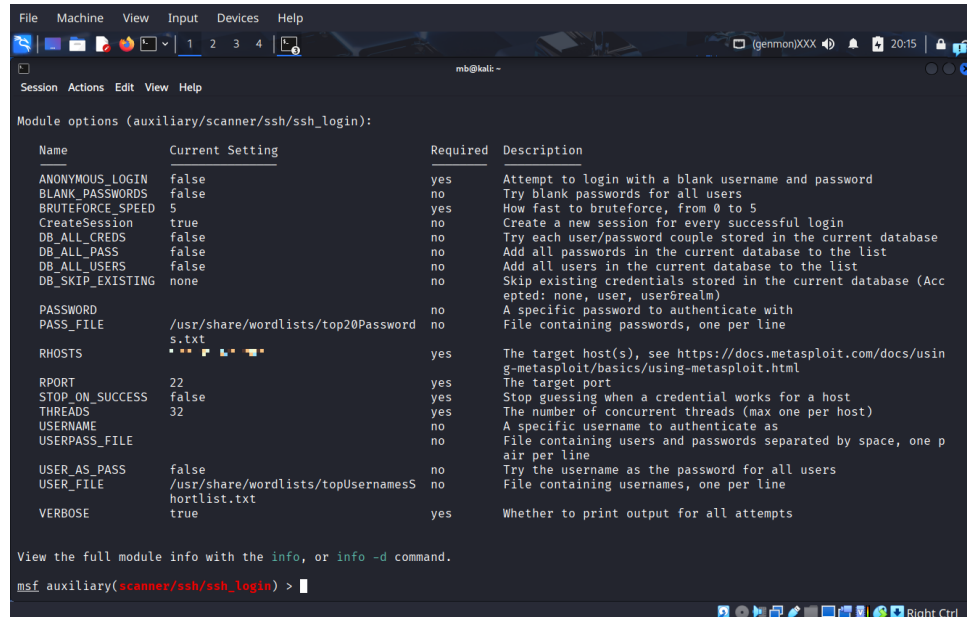


Figure 4 – Starting up Metasploit

Afterward, I switched to the `ssh_login` tool and configured the options by setting the host IP, the path to the password list file, the path to the username list file, the number of threads to 32, and the verbosity to true. To expedite the experiment, I created custom passwords and username wordlists, containing only the most popular credentials.



```
File  Machine  View  Input  Devices  Help
1  2  3  4
mb@kali -

Session  Actions  Edit  View  Help

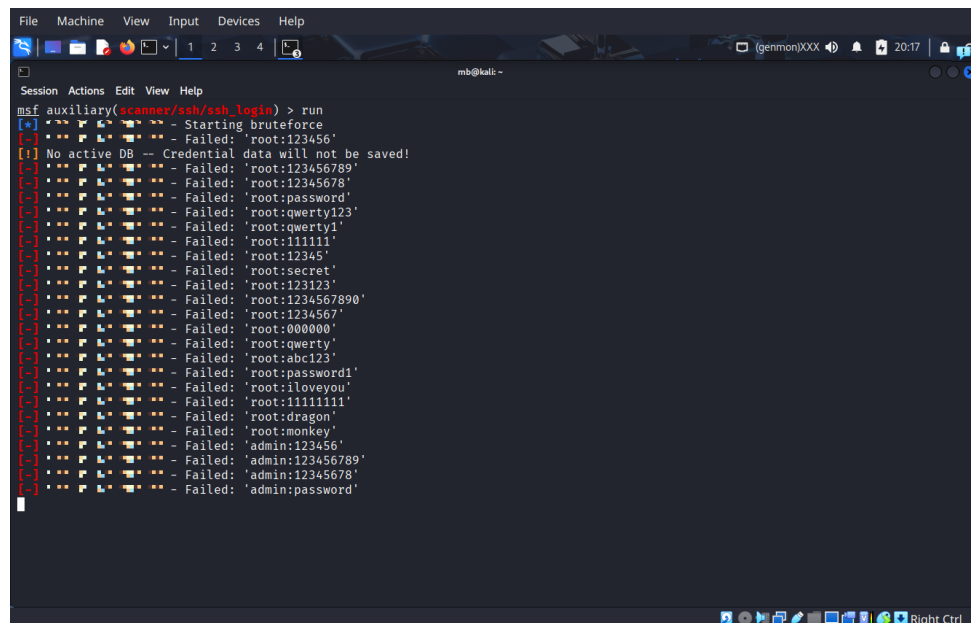
Module options (auxiliary/scanner/ssh/ssh_login):

  Name                Current Setting      Required  Description
  ----                -
  ANONYMOUS_LOGIN      false                yes       Attempt to login with a blank username and password
  BLANK_PASSWORDS      false                no        Try blank passwords for all users
  BRUTEFORCE_SPEED     5                    yes       How fast to bruteforce, from 0 to 5
  CreateSession        true                 no        Create a new session for every successful login
  DB_ALL_CREDS         false                no        Try each user/password couple stored in the current database
  DB_ALL_PASS          false                no        Add all passwords in the current database to the list
  DB_ALL_USERS         false                no        Add all users in the current database to the list
  DB_SKIP_EXISTING     none                 no        Skip existing credentials stored in the current database (Accepted: none, user, user@realm)
  PASSWORD             no                   no        A specific password to authenticate with
  PASS_FILE            /usr/share/wordlists/top20Password s.txt               no        File containing passwords, one per line
  RHOSTS               *.*.*.*.*            yes       The target host(s), see https://docs.metasploit.com/docs/using-metasploit/basics/using-metasploit.html
  RPORT               22                   yes       The target port
  STOP_ON_SUCCESS      false                yes       Stop guessing when a credential works for a host
  THREADS              32                   yes       The number of concurrent threads (max one per host)
  USERNAME             no                   no        A specific username to authenticate as
  USERPASS_FILE       no                   no        File containing users and passwords separated by space, one pair per line
  USER_AS_PASS         false                no        Try the username as the password for all users
  USER_FILE            /usr/share/wordlists/topUsernamesS hortlist.txt         no        File containing usernames, one per line
  VERBOSE              true                 yes       Whether to print output for all attempts

View the full module info with the info, or info -d command.
msf auxiliary(scanner/ssh/ssh_login) > |
```

Figure 5 – Metasploit `ssh_login` options

I started the attack, and it began passing the passwords and usernames. However, the attack was notably slow, even with a high thread count and sufficient computing power allocated to the VM. After research, I determined this was due to SSH's built-in **brute force mitigations**, which limit login retries. Despite knowing these initial tests would not successfully gain access (as I intentionally withheld the correct credentials), I proceeded with implementing additional security controls to further harden the system.



```
msf auxiliary(scanner/ssh_login) > run
[*] - Starting brute force
[*] - Failed: 'root:123456'
[*] - Failed: 'root:123456789'
[*] - Failed: 'root:12345678'
[*] - Failed: 'root:password'
[*] - Failed: 'root:qwerty123'
[*] - Failed: 'root:qwerty1'
[*] - Failed: 'root:1111111'
[*] - Failed: 'root:12345'
[*] - Failed: 'root:secret'
[*] - Failed: 'root:123123'
[*] - Failed: 'root:1234567890'
[*] - Failed: 'root:1234567'
[*] - Failed: 'root:000000'
[*] - Failed: 'root:qwerty'
[*] - Failed: 'root:abc123'
[*] - Failed: 'root:password1'
[*] - Failed: 'root:iloveyou'
[*] - Failed: 'root:11111111'
[*] - Failed: 'root:dragon'
[*] - Failed: 'root:monkey'
[*] - Failed: 'admin:123456'
[*] - Failed: 'admin:123456789'
[*] - Failed: 'admin:12345678'
[*] - Failed: 'admin:password'
```

Figure 6 – running the brute force attack on the Kali Linux VM

Now that the initial test attack had finished, it was time to harden the system using Iptables and Splunk.

Phase 2: Defense

Step 1: setup Iptable rules

Switching to the host machine, I opened the terminal and added the following iptables rules:

```
sudo iptables -I INPUT -p tcp --dport 22 -m state --state NEW -m recent --set
```

```
sudo iptables -I INPUT -p tcp --dport 22 -m recent -update -seconds 60 --hitcount 4 -j DROP
```

I used Gemini 2.5 Flash to learn the Iptables syntax and the meaning of each command component. I utilized the `-I` flag to ensure the rules were inserted at the top of the rule list, which prioritizes them based on how firewall rules are processed. Below is a snapshot of my terminal entering these Iptables rules.



Figure 7 – adding Iptable rules

To check if the rules were added, I used the following command:

```
sudo iptables -L INPUT --line-numbers
```

The output is below.

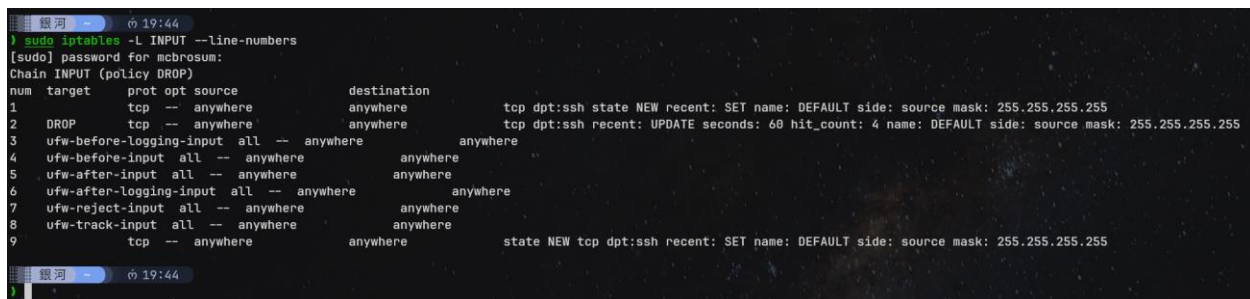


Figure 8 – Iptable output in the terminal

Next is Splunk monitoring.

Step 2: Add authentication logs into Splunk

I had prior experience with Splunk, so I was familiar with the basics of the UI; however, I needed to study the search syntax more deeply. Gemini 2.5 Flash was helpful for learning how to search logs, and I also followed a specific Splunk guide for searching SSH logs [1]. Below is my first search for invalid SSH login attempts:

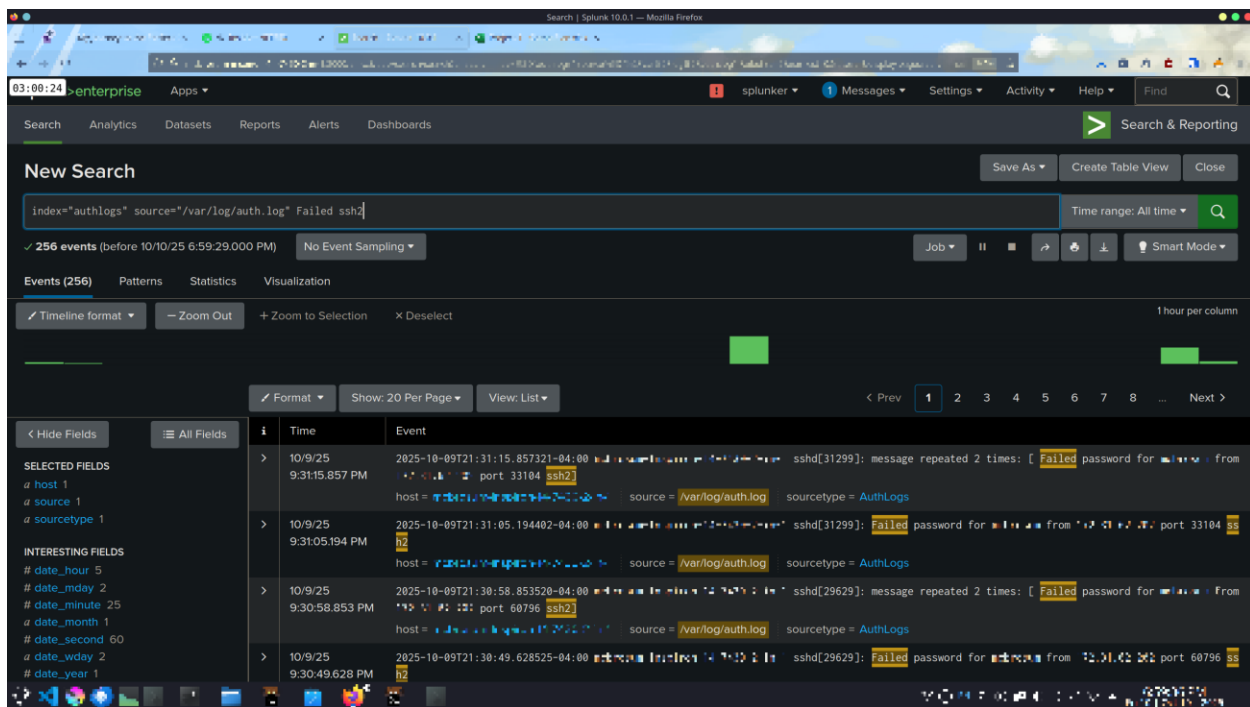


Figure 9 – Splunk search results with “Failed” and “ssh2” keywords.

After finding the correct logs, I extracted a field and named it SSHInvalidUser by using regex. Adding this field to the search result creates this chart:

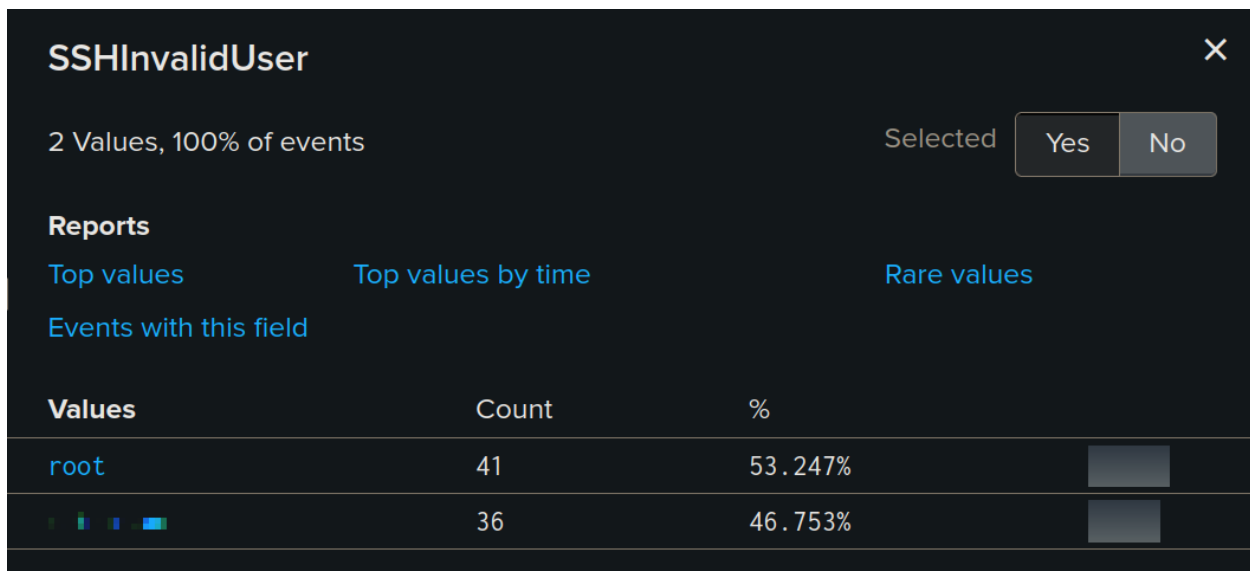
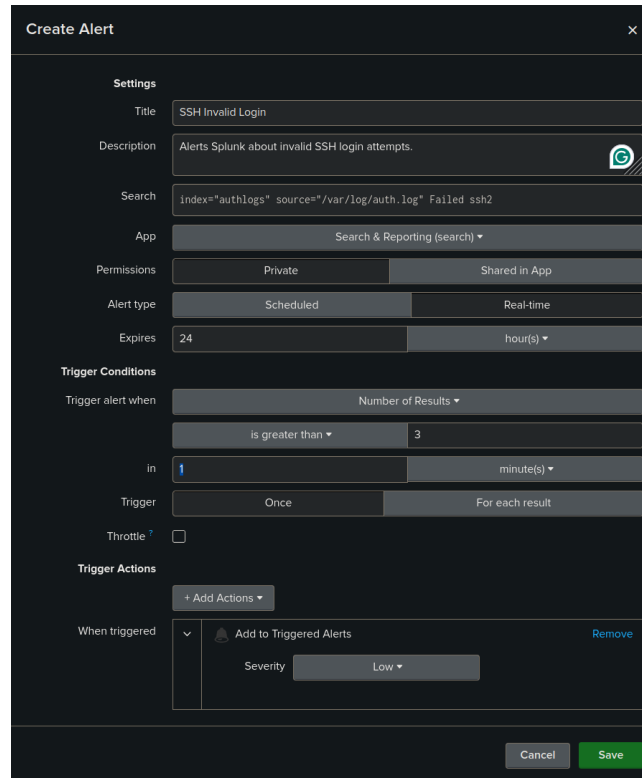


Figure 10 – chart showing invalid users from the SSH attack.

Step 3: Create Splunk alerts for SSH login attempts

Once I found the correct logs, I created real-time alerts using the SSH log search query. Below is a depiction of me setting these alerts up:



The screenshot shows the 'Create Alert' interface in Splunk. The 'Settings' section includes: Title 'SSH Invalid Login', Description 'Alerts Splunk about invalid SSH login attempts.', Search query 'index="authlogs" source="/var/log/auth.log" Failed ssh2', App 'Search & Reporting (search)', Permissions 'Private', Alert type 'Real-time', and Expires '24 hour(s)'. The 'Trigger Conditions' section shows 'Trigger alert when' set to 'Number of Results' 'is greater than' '3' 'in' '1 minute(s)'. The 'Trigger' is set to 'Once'. The 'Throttle' checkbox is unchecked. The 'Trigger Actions' section shows 'When triggered' with a dropdown menu containing 'Add to Triggered Alerts' and 'Remove'. The 'Severity' is set to 'Low'. The 'Cancel' and 'Save' buttons are at the bottom right.

Figure 10 – setting up real-time SSH Splunk alerts.

Now that my defenses are set up, I can initiate a test attack.

Phase 3: Test Attack

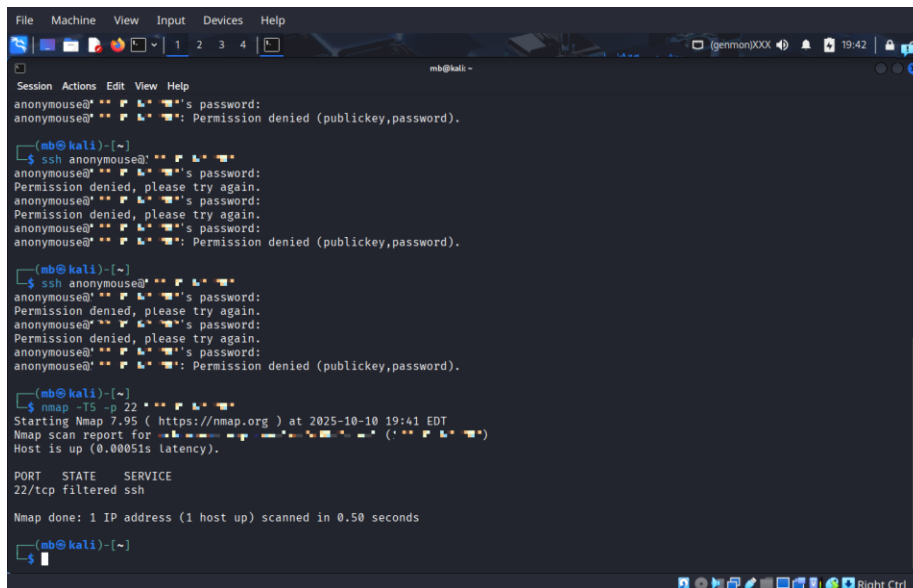
Step 1: Brute force port 22

Using the same Kali Linux virtual machine and method as before, I attempted the same attack on the target machine. This time, I manually inputted random passwords instead of using Metasploit. This was done specifically to trigger the Splunk alerts and verify that the firewall dropped my packets as intended.

Step 2: View port 22 with Nmap after attack

Going back to the attack machine, I ran another Nmap scan after the attack. The port state is now reported as "filtered," meaning that packets from the attack machine are being

dropped before reaching the target machine. The attack machine will now have to wait one minute before it can make four more attempts.

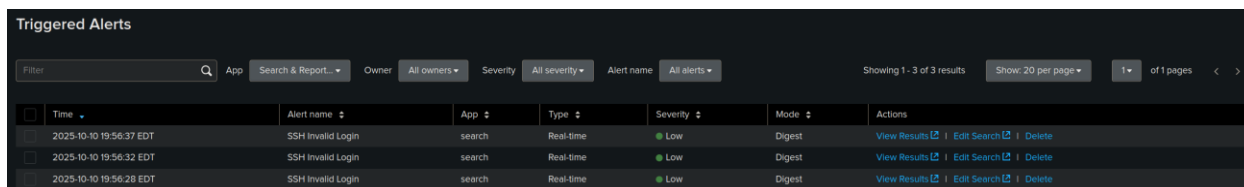


```
mb@kali: ~  
$ ssh anonymouse@10.10.10.22  
anonymouse@10.10.10.22: Permission denied (publickey,password).  
  
mb@kali: ~  
$ ssh anonymouse@10.10.10.22  
anonymouse@10.10.10.22: Permission denied (publickey,password).  
anonymouse@10.10.10.22: Permission denied, please try again.  
anonymouse@10.10.10.22: Permission denied, please try again.  
anonymouse@10.10.10.22: Permission denied, please try again.  
anonymouse@10.10.10.22: Permission denied (publickey,password).  
  
mb@kali: ~  
$ ssh anonymouse@10.10.10.22  
anonymouse@10.10.10.22: Permission denied (publickey,password).  
anonymouse@10.10.10.22: Permission denied, please try again.  
anonymouse@10.10.10.22: Permission denied, please try again.  
anonymouse@10.10.10.22: Permission denied, please try again.  
anonymouse@10.10.10.22: Permission denied (publickey,password).  
  
mb@kali: ~  
$ nmap 10.10.10.22  
Starting Nmap 7.95 ( https://nmap.org ) at 2025-10-10 19:41 EDT  
Nmap scan report for 10.10.10.22: 22/tcp filtered ssh  
Host is up (0.00051s latency).  
  
PORT      STATE SERVICE  
22/tcp    filtered ssh  
  
Nmap done: 1 IP address (1 host up) scanned in 0.50 seconds  
  
mb@kali: ~  
$
```

Figure 12 – An Nmap scan showing a filtered state on port 22.

Step 2: Monitor SSH login alerts in Splunk

While the attack was live, I observed alerts on Splunk. Splunk successfully generated alerts based on the failed SSH logins.



Triggered Alerts						
Filter	App	Search & Report...	Owner	Severity	Alert name	All alerts
Showing 1 - 3 of 3 results Show: 20 per page 1 of 1 pages						
<input type="checkbox"/>	Time	Alert name	App	Type	Severity	Mode
<input type="checkbox"/>	2025-10-10 19:56:37 EDT	SSH Invalid Login	search	Real-time	Low	Digest
<input type="checkbox"/>	2025-10-10 19:56:32 EDT	SSH Invalid Login	search	Real-time	Low	Digest
<input type="checkbox"/>	2025-10-10 19:56:28 EDT	SSH Invalid Login	search	Real-time	Low	Digest

Figure 11 – A view of custom Splunk alerts.

This concludes the experiment.

Outcome

As my first project, testing and hardening my machine was an enjoyable and valuable exercise. I gained significant knowledge and look forward to performing more such experiments in the future.

Lessons Learned

The following is a summary of what I learned in this project:

1. Splunk
 - a. How to add data specify which files to monitor
 - b. How to search for logs using regex, specifying the source
 - c. How to extract fields from logs
 - d. How to create and view alerts
2. SSH
 - a. SSH's built in brute force protection
 - b. How to download and start SSH
3. Virtual Machine
 - a. How the virtual machine gets Wi-Fi connection implicitly through the host
4. Iptables
 - a. Overall syntax
 - b. Listing rules
 - c. Creating new rules

I used online resources and Gemini Pro to learn and practice for this project.

Recommendations

For further hardening of SSH connections, you can disable root login and require Multi-Factor Authentication (MFA). Disabling root login prevents attackers from immediately gaining a privileged shell, and MFA provides additional mitigation against brute-force attacks.

For future projects, I will increase my use of AI for learning tool usage and rely more on online video tutorials for implementing security controls. Reading numerous articles that did not directly apply to the specific situation extended the overall project work time.

References

[1]

D. Miessler, "Monitoring SSH Bruteforce Attempts Using Splunk," *Danielmiessler.com*, 2017. <https://danielmiessler.com/blog/monitoring-ssh-bruteforce-attempts-using-splunk> (accessed Oct. 11, 2025).